

# RPCS AND GRPC

George Porter  
Apr 24, 2018

UC San Diego



## ATTRIBUTION

- These slides are released under an Attribution-NonCommercial-ShareAlike 3.0 Unported (CC BY-NC-SA 3.0) Creative Commons license
- These slides incorporate material from:
  - Kyle Jamieson, Princeton University (also under a CC BY-NC-SA 3.0 Creative Commons license)

UC San Diego



## Outline

1. RPC fundamentals
2. Handling failures in RPCs
3. gRPC: Google RPC overview
4. ATM Server demo

## WHY RPC?

- The typical programmer is trained to write single-threaded code that runs in **one place**
- **Goal:** Easy-to-program network communication that makes client-server communication **transparent**
  - Retains the “feel” of writing centralized code
    - Programmer needn’t think about the network

## REMOTE PROCEDURE CALL (RPC)

- Distributed programming is challenging
  - Need common primitives/abstraction to hide complexity
  - E.g., file system abstraction to hide block layout, process abstraction for scheduling/fault isolation
- In early 1980's, researchers at PARC noticed most distributed programming took form of *remote procedure call*

## WHAT'S THE GOAL OF RPC?

- Within a single program, running in a single process, recall the well-known notion of a procedure call:
  - Caller pushes arguments onto stack,
    - jumps to address of callee function
  - Callee reads arguments from stack,
    - executes, puts return value in register,
    - returns to next instruction in caller

RPC's Goal: To make communication appear like a local procedure call: transparency for procedure calls

## RPC EXAMPLE

### Local computing

```
X = 3 * 10;
print(X)
> 30
```

### Remote computing

```
server = connectToServer(S);
Try:
    X = server.mult(3,10);
    print(X)
Except e:
    print "Error!"
> 30
or
> Error
```

## RPC ISSUES

- Heterogeneity
  - Client needs to **rendezvous** with the server
  - Server must **dispatch** to the required function
    - What if server is **different** type of machine?
- Failure
  - What if messages get **dropped**?
  - What if client, server, or network **fails**?
- Performance
  - Procedure call takes  $\approx 10$  cycles  $\approx 3$  ns
  - RPC in a data center takes  $\approx 10$   $\mu$ s ( $10^3\times$  slower)
    - In the wide area, typically  $10^6\times$  slower

## PROBLEM: DIFFERENCES IN DATA REPRESENTATION

- Not an issue for **local** procedure call
- For a remote procedure call, a **remote machine may**:
  - Represent data types using **different sizes**
  - Use a **different byte ordering** (*endianness*)
  - Represent floating point numbers **differently**
  - Have **different data alignment** requirements
    - e.g., 4-byte type begins only on 4-byte memory boundary

## BYTE ORDER

- x86-64 is a **little endian** architecture

- **Least** significant byte of multi-byte entity at **lowest** memory address

- “Little end goes first”

- Some other systems use **big endian**

- **Most** significant byte of multi-byte entity at **lowest** memory address

- “Big end goes first”

int 5 at address 0x1000:

0x1000:	0000 0101
0x1001:	0000 0000
0x1002:	0000 0000
0x1003:	0000 0000

int 5 at address 0x1000:

0x1000:	0000 0000
0x1001:	0000 0000
0x1002:	0000 0000
0x1003:	0000 0101

## PROBLEM: DIFFERENCES IN PROGRAMMING SUPPORT

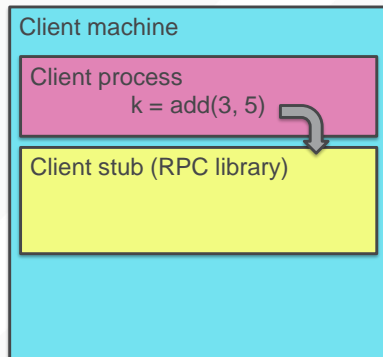
- Language support **varies**:
  - Many programming languages have **no inbuilt concept** of remote procedure calls
    - *e.g.*, C, C++, earlier Java
  - Some languages have **support that enables RPC**
    - *e.g.*, Python, Haskell, Go

## SOLUTION: INTERFACE DESCRIPTION LANGUAGE

- Mechanism to pass procedure parameters and return values in a **machine-independent way**
- Programmer may write an **interface description** in the IDL
  - Defines API for procedure calls: names, parameter/return types
- Then runs an **IDL compiler** which generates:
  - Code to **marshal** (convert) native data types into machine-independent byte streams
    - And vice-versa, called **unmarshaling**
  - **Client stub**: Forwards local procedure call as a request to server
  - **Server stub**: Dispatches RPC to its implementation

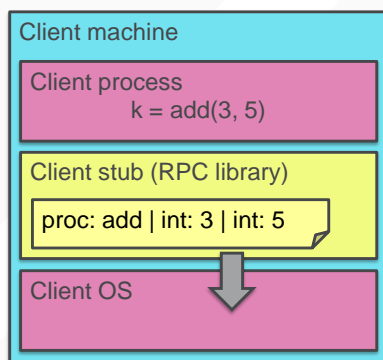
## A DAY IN THE LIFE OF AN RPC

1. Client calls stub function (pushes params onto stack)




## A DAY IN THE LIFE OF AN RPC

1. Client calls stub function (pushes params onto stack)
2. Stub marshals parameters to a network message



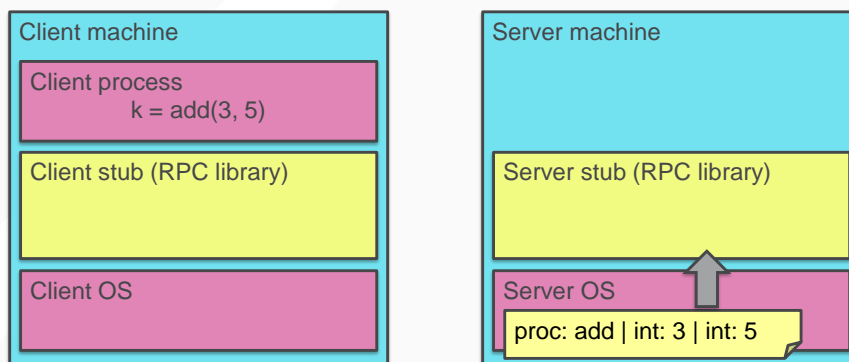
## A DAY IN THE LIFE OF AN RPC

2. Stub marshals parameters to a network message
- 3. OS sends a network message to the server**



## A DAY IN THE LIFE OF AN RPC

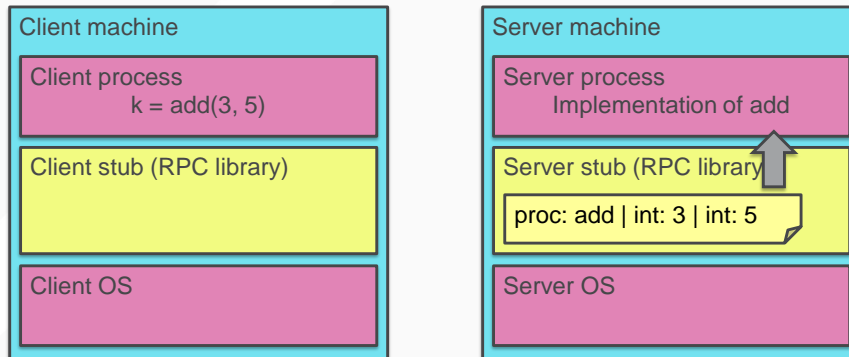
3. OS sends a network message to the server
- 4. Server OS receives message, sends it up to stub**





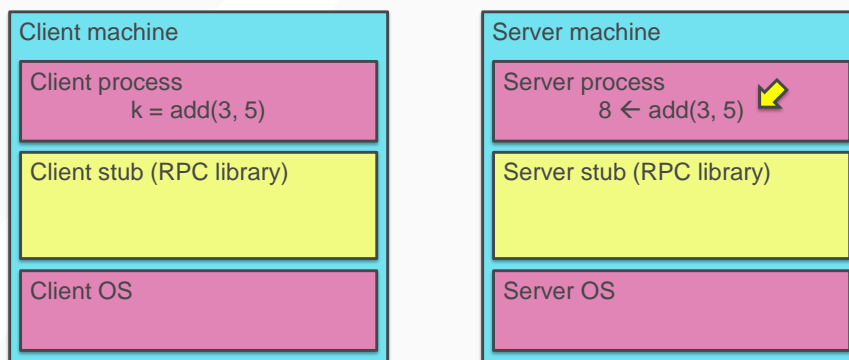
## A DAY IN THE LIFE OF AN RPC

4. Server OS receives message, sends it up to stub
5. **Server stub unmarshals params, calls server function**



## A DAY IN THE LIFE OF AN RPC

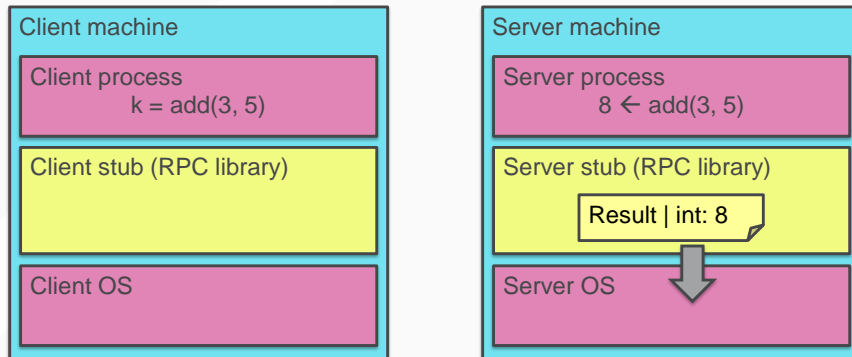
5. Server stub unmarshals params, calls server function
6. **Server function runs, returns a value**



## A DAY IN THE LIFE OF AN RPC

6. Server function runs, returns a value

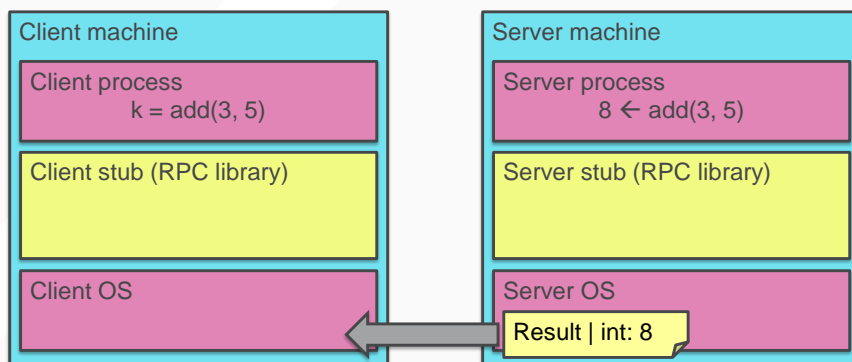
**7. Server stub marshals the return value, sends msg**



## A DAY IN THE LIFE OF AN RPC

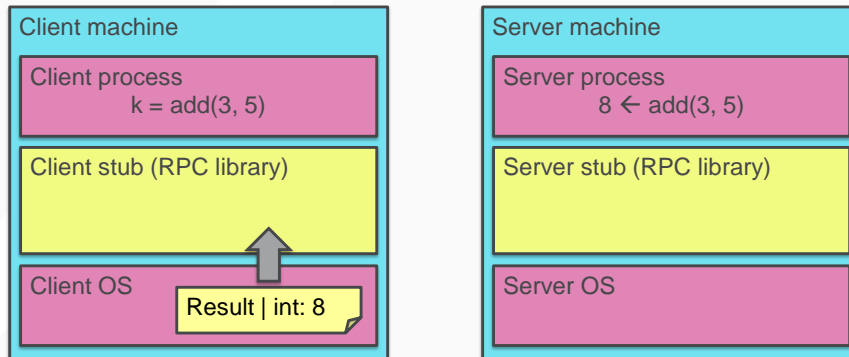
7. Server stub marshals the return value, sends msg

**8. Server OS sends the reply back across the network**



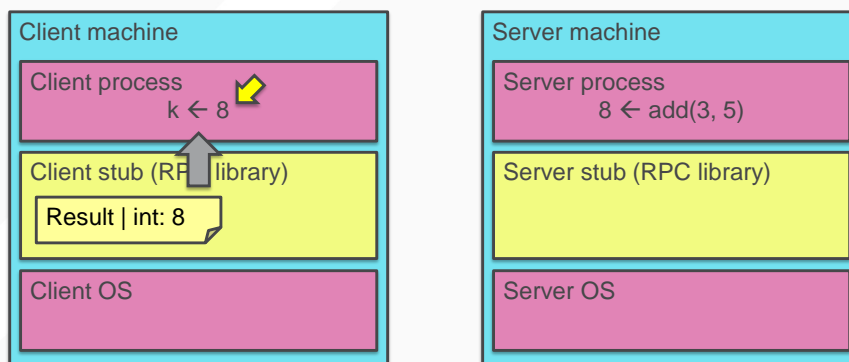
## A DAY IN THE LIFE OF AN RPC

8. Server OS sends the reply back across the network
- 9. Client OS receives the reply and passes up to stub**




## A DAY IN THE LIFE OF AN RPC

9. Client OS receives the reply and passes up to stub
- 10. Client stub unmarshals return value, returns to client**



## PETERSON AND DAVIE VIEW



## THE SERVER STUB IS REALLY TWO PARTS

- **Dispatcher**
  - Receives a client's RPC request
    - **Identifies** appropriate server-side method to invoke
- **Skeleton**
  - **Unmarshals** parameters to server-native types
  - **Calls** the local server procedure
  - **Marshals** the response, sends it back to the dispatcher
- **All this is hidden from the programmer**
  - Dispatcher and skeleton may be integrated
    - Depends on implementation



## Outline

1. RPC fundamentals
2. Handling failures in RPCs
3. gRPC: Google RPC overview
4. ATM Server demo

## WHAT COULD POSSIBLY GO WRONG?

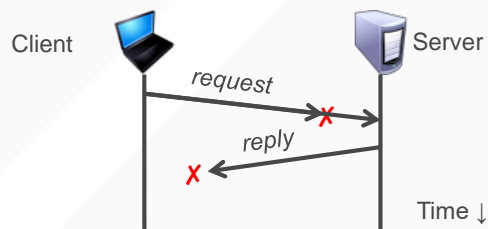


## WHAT COULD POSSIBLY GO WRONG?

1. Client may **crash and reboot**
2. Packets may be **dropped**
  - Some individual **packet loss** in the Internet
  - **Broken routing** results in many lost packets
3. Server may **crash and reboot**
4. Network or server might just be **very slow**

All these may **look the same** to the client...

## FAILURES, FROM CLIENT'S PERSPECTIVE




The cause of the failure is **hidden** from the **client**!

## AT-LEAST-ONCE SCHEME

- **Simplest** scheme for handling failures
- 1. Client stub **waits for a response**, for a while
  - Response takes the form of an **acknowledgement** message from the server stub
- 2. If no response arrives after a fixed **timeout** time period, then client stub **re-sends the request**
- Repeat the above a few times
  - *Still no response?* Return an error to the application

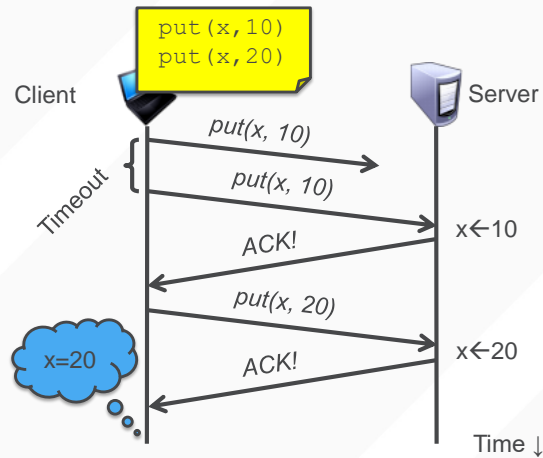
## AT-LEAST-ONCE AND SIDE EFFECTS

- Client sends a “debit \$10 from bank account” RPC




## AT-LEAST-ONCE AND WRITES

- $\text{put}(x, \text{value})$ , then  $\text{get}(x)$ : expect answer to be *value*



## AT-LEAST-ONCE AND WRITES

- Consider a client storing **key-value pairs** in a **database**
- $\text{put}(x, \text{value})$ , then  $\text{get}(x)$ : expect answer to be *value*





## SO IS AT-LEAST-ONCE *EVER* OKAY?

- **Yes:** If they are read-only operations with no side effects
  - *e.g.*, read a key's value in a database
- **Yes:** If the application has its own functionality to cope with duplication and reordering
  - You will implement this in Project 2

## AT-MOST-ONCE SCHEME

- **Idea:** server RPC code detects duplicate requests
  - Returns previous reply **instead of re-running handler**
- *How to detect a duplicate request?*
  - **Test:** Server sees same function, same arguments twice
    - **No!** Sometimes applications **legitimately** submit the same function with same arguments, twice in a row

## AT-MOST-ONCE SCHEME

- *How to detect a duplicate request?*
- Client includes unique **transaction ID (xid)** with each one of its RPC requests
- Client uses **same xid** for retransmitted requests

```

At-Most-Once Server
if seen[xid]:
    retval = old[xid]
else:
    retval = handler()
    old[xid] = retval
    seen[xid] = true
return retval
  
```

## AT MOST ONCE: ENSURING UNIQUE XIDS

- *How to ensure that the xid is unique?*
1. Combine a unique client ID (e.g., IP address) with the current time of day
  2. Combine unique client ID with a sequence number
    - Suppose the client crashes and restarts. *Can it reuse the same client ID?*
  3. Big random number

## AT-MOST-ONCE: DISCARDING SERVER STATE

- **Problem:** `seen` and `old` arrays will **grow without bound**
- **Observation:** By construction, when the client gets a response to a particular `xid`, it will **never re-send it**
- Client could **tell** server “I’m done with `xid x` – delete it”
  - Have to tell the server about **each and every** retired `xid`
    - Could **piggyback** on subsequent requests

Significant overhead if many RPCs are in flight, in parallel

## AT-MOST-ONCE: DISCARDING SERVER STATE

- **Problem:** `seen` and `old` arrays will **grow without bound**
- Suppose `xid` =  $\langle \text{unique client id, sequence no.} \rangle$ 
  - e.g.  $\langle 42, 1000 \rangle, \langle 42, 1001 \rangle, \langle 42, 1002 \rangle$
- Client includes “seen all replies  $\leq X$ ” with every RPC
  - Much like TCP sequence numbers, acks
- *How does the client **know** that the server received the information about retired RPCs?*
  - Each one of these is cumulative: later seen messages subsume earlier ones

## AT-MOST-ONCE: CONCURRENT REQUESTS

- **Problem:** How to handle a duplicate request while the original is still executing?
  - Server doesn't know reply yet. Also, we don't want to run the procedure twice
- **Idea:** Add a **pending** flag per executing RPC
  - Server waits for the procedure to finish, or ignores

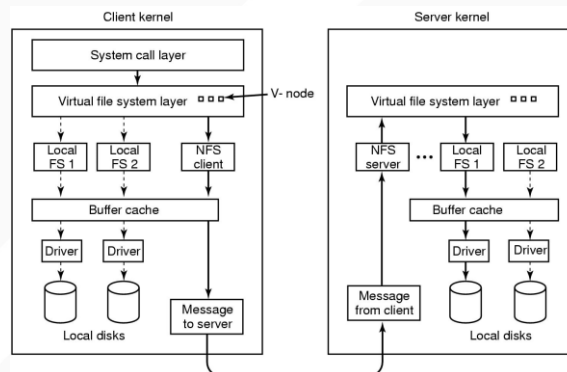
## AT MOST ONCE: SERVER CRASH AND RESTART

- **Problem:** Server may crash and restart
- *Does server need to write its tables to disk?*
- Yes! On **server crash and restart:**
  - If `old[]`, `seen[]` tables are only in memory:
    - Server will forget, **accept duplicate requests**

## RPC SEMANTICS

Delivery Guarantees			RPC Call Semantics
Retry Request	Duplicate Filtering	Retransmit Response	
No	NA	NA	<i>Maybe</i>
Yes	No	Re-execute Procedure	<i>At-least once</i>
Yes	Yes	Retransmit reply	<i>At-most once</i>

## SUMMARY: RPC



- RPC everywhere!
- **Necessary** issues surrounding machine heterogeneity
- **Subtle** issues around handling **failures**



## Outline

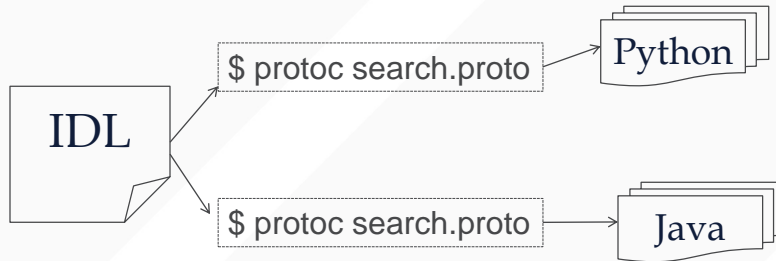
1. RPC fundamentals
2. Handling failures in RPCs
3. gRPC: Google RPC overview
4. ATM Server demo

## GOOGLE RPC (GRPC)

- Cross-platform RPC toolkit developed by Google
- Languages:
  - C++, Java, Python, Go, Ruby, C#, Node.js, Android, Obj-C, PHP
- Defines *services*
  - Collection of RPC calls

```
service Search {  
  rpc searchWeb(SearchRequest) returns (SearchResult) {}  
}
```

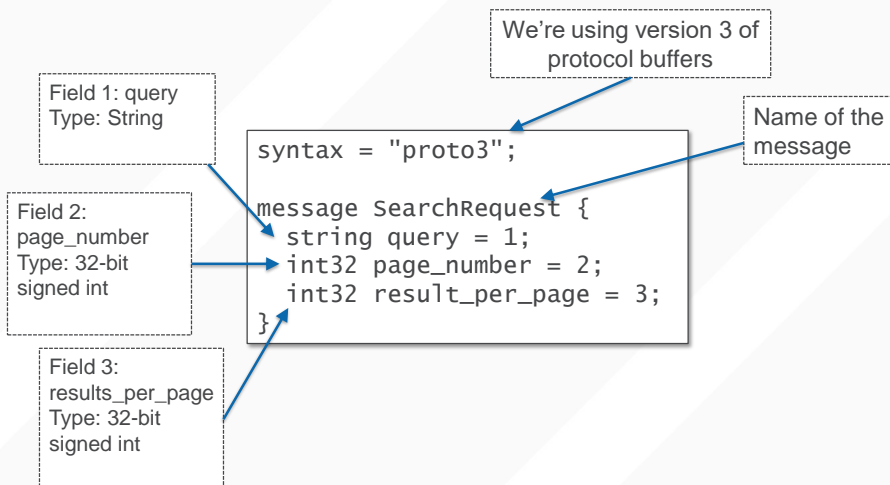
## IDL: INTERFACE DEFINITION LANGUAGE



- Language-neutral way of specifying:
  - Data structures (called Messages)
  - Services, consisting of procedures/methods
- Stub compiler
  - Compiles IDL into Python, Java, etc.

## IDL LANGUAGE: PROTOCOL BUFFERS

- Defines Messages (i.e., data structures)



## PROTOCOL BUFFERS: BASE TYPES

• <u>protobuf IDL:</u>	• <u>Python:</u>	• <u>Java:</u>	• <u>C++:</u>
• double, float	• float, float	• double, float	• double, float
• int32, int64	• int, int/long	• int, long	• int32, int64
• uint32, uint64	• int, int/long	• int, long	• uint32, uint64
• bool	• bool	• Boolean	• bool
• string	• str	• String	• string
• bytes	• str	• ByteString	• string

## IDL POSITIONAL ARGUMENTS

- Why do we label the fields with numbers?
- So we can change “signature” of the message later and **still be compatible** with legacy code

```
syntax = "proto3";

message SearchRequest {
  string query = 1;
  int32 page_number = 2;
  int32 result_per_page = 3;
}
```



```
syntax = "proto3";

message SearchRequest {
  string query = 1;
  int32 page_number = 2;
  int32 shard_num = 4;
}
```



## MAKING SERVICES *EVOLVABLE*

- No way to “stop everything” and upgrade
- Clients/servers/services must co-exist
- For newly added fields, old services use defaults:
  - String: ""
  - bytes: []
  - bools: false
  - numeric: 0
  - ...

## PROTOCOL BUFFERS: MAP TYPE

- `map<key_type, value_type> map_field = N;`
- Example:
  - `map<string, Project> projects = 3;`

## IMPLEMENTING IN DIFFERENT LANGUAGES

### IDL

```
message Person {
  required string name = 1;
  required int32 id = 2;
  optional string email = 3;
}
```

### C++: reading from a file

```
Person john;
fstream input(argv[1],
              ios::in | ios::binary);
john.ParseFromIstream(&input);
id = john.id();
name = john.name();
email = john.email();
```

### Java: writing to a file

```
Person john = Person.newBuilder()
    .setId(1234)
    .setName("John Doe")
    .setEmail("jdoe@example.com")
    .build();
output = new FileOutputStream(args[0]);
john.writeTo(output);
```

## A C++ EXAMPLE

```
Person person;
person.set_name("John Doe");
person.set_id(1234);
person.set_email("jdoe@example.com");
fstream output("myfile", ios::out | ios::binary);
person.SerializeToOstream(&output);
```

```
fstream input("myfile", ios::in | ios::binary);
Person person;
person.ParseFromIstream(&input);
cout << "Name: " << person.name() << endl;
cout << "E-mail: " << person.email() << endl;
```

- Can read/write protobuf Message objects to files/stream/raw sockets
- In particular, gRPC service RPCs
  - Take Message as argument, return Message as response

## JAVA OVERVIEW: THE SERVER

- Outer class (e.g., `ATMServer`)
  - `start()`: starts the server (listens for incoming connections)
  - `stop()`: stops the server
  - `blockUntilShutdown()`: internal method
  - `main()`: reads command-line arguments, calls `start()`
  - static class `ATMServerImpl` extends `ATMServerGrpc...`
    - This is where you specify the implementation of the RPCs

## ATMSERVERIMPL

```
message SearchRequest {
    string query = 1;
}

message SearchResponse {
    int32 numResponses = 1;
    repeated string response = 2;
}

service Search {
    rpc searchWeb(SearchRequest) returns (SearchResult) {}
}
```

- Implement each RPC here by overriding the compiler-generated code
- Compiler output goes in `target/generated-sources` folder
  - `Generated-sources/java` contains 'message' definitions
  - `Generated-sources/grpc-java` contains 'service' definitions

## HOW TO GET STARTED?

1. Define your messages and services in proto file
2. Run the compiler
 

```
$ mvn protobuf:compile protobuf:compile-custom
```
3. Find proc signature in .../ATMServerGrpc.java

```
/**
 */
public static abstract class ATMServerImplBase implements io.grpc.BindableService {

    /**
     */
    public void ping(atm.PingRequest request,
        io.grpc.stub.StreamObserver<atm.PingResponse> responseObserver) {
        asyncUnaryCall(METHOD_PING, responseObserver);
    }

    /**
     */
    public void searchWeb(atm.SearchRequest request,
        io.grpc.stub.StreamObserver<atm.SearchResponse> responseObserver) {
        asyncUnaryCall(METHOD_SEARCH_WEB, responseObserver);
    }
}
```

4. Override in src/main/java/atm/ATMServer.java

## UNDERSTANDING STUB CODE

```
service Search {
  rpc searchWeb(SearchRequest) returns
    (SearchResult) {}
}
```

- Yet stub code looks like this:

```
/**
 */
public static abstract class ATMServerImplBase implements io.grpc.BindableService {

    /**
     */
    public void ping(atm.PingRequest request,
        io.grpc.stub.StreamObserver<atm.PingResponse> responseObserver) {
        asyncUnaryCall(METHOD_PING, responseObserver);
    }

    /**
     */
    public void searchWeb(atm.SearchRequest request,
        io.grpc.stub.StreamObserver<atm.SearchResponse> responseObserver) {
        asyncUnaryCall(METHOD_SEARCH_WEB, responseObserver);
    }
}
```

- Void return type
- Always two arguments
  - 1<sup>st</sup> argument is a single message which is the only argument the RPC can take
    - SearchRequest in this case
  - 2<sup>nd</sup> argument is an “in-out” argument used to return data back to the caller
- gRPC defines four types of in-out return types—we’re only using the most basic one
- To return data to client, you pass it to the “responseObserver”

## CONSTRUCTING RESULTS FROM STUB CODE

1. Construct an object based on the return type

• Example:

- `SearchResultBuilder srb;`
- `srb.setQuery("foo bar");`
- `SearchResult sr = srb.build();`

## RETURNING RESULTS FROM STUB CODE

1. Construct an object based on the return type

2. Pass to the responseObserver

3. Tell the responseObserver you're done

• Example:

- `responseObserver.onNext(sr);`
- `responseObserver.onCompleted();`

## ERRORS AND EXCEPTIONS

- The server can throw an Exception, which is translated into an Exception in the client
  - Catch try...catch and handle as appropriate

## DIVING DEEPER

- <https://grpc.io/>
- <https://grpc.io/docs/quickstart/java.html>
- <https://grpc.io/docs/tutorials/basic/java.html>
- <https://grpc.io/docs/reference/java/generated-code.html>



## Outline

1. RPC fundamentals
2. Handling failures in RPCs
3. gRPC: Google RPC overview
4. ATM Server demo

## SIMPLE ATM SERVER



- Operations:
  - login
    - Account number + PIN
  - deposit
    - \$\$\$
  - getBalance
  - logout

## SIMPLE ATM SERVER



- Keeping track of account + pin with “login tokens”
- After logging in, get a token
- Use token to deposit money, withdraw, transfer, ...

UC San Diego